

Homework requirements

You are required to turn in two files `recommend.py`, containing the required functions, and `README.txt`, which must specify which level you are attempting/claiming to have completed. As always, if you do the homework with a partner, both of you must include your names in the `README.txt` file.

To help you gauge your progress and test your code, we will develop a test suite you can use to test your code. Don't treat the test suite as complete: we will test and inspect your code using more tests than those we have provided.

Grade level descriptions

C-level

You will write two functions:

- `readRatings(ratings_file)`
 - Reads in a ratings file in the format described above. The input is the location of the ratings file (i.e., `readRatings("ratings.tsv")`). The output is a dictionary whose keys are `user_ids`, and whose values are dictionaries whose keys are movie titles and whose values are the user's rating for those movies. For example:
 - ```
{
 'Terveen': {
 'Pulp Fiction': 5,
 'Armageddon': 2.0,
 'Star Wars: Episode I - The Phantom Menace': 1.0,
 'Once Upon a Time in Mexico': 3.0
 },
 'Kluver': {
 'Harold and Kumar Go To White Castle': 5.0,
 'Bridges of Madison County': 5.0,
 'Lord of the Rings: Return of The King': 2.0,
 'BattleField Earth': 4.0
 }
}
```
  - Using the `split()` function will make it very easy to separate the `user_id`, `rating`, and `movie_title` fields
- `similarity(user_ratings_1, user_ratings_2)`
  - `user_ratings_1` and `user_ratings_2` are dictionaries that contain the ratings for two users in the format produced by `readRatings()`.

- This function returns a float value between 1 (indicating total agreement) and -1 (indicating total disagreement).

You will compute similarity using the Pearson correlation coefficient. This is described in the [Wikipedia article on collaborative filtering](#): since this is stated in terms of user ratings, you can follow it pretty exactly in your implementation.

**IMPORTANT CLARIFICATION:** When you compute the term for the average of each user's ratings, you should include **all** the user's ratings, not just ratings of items in the intersection of the two user's ratings.

Because some people in the class have seen different movies you should perform the computation only over the movies that both users have rated. If the users have not rated any movies in common you should return a similarity of 0.

As always, you may define any helper functions you find useful, for this and any of the levels of the assignment.

## B-Level

You will implement the k-nearest-neighbors algorithm using your similarity function.

- `nearestNeighbors(user_id, all_user_ratings, k)`
  - This function returns a list of the k "nearest neighbors" of the user `user_id`. `k` is an integer (don't worry about error checking -- you can assume it really is an integer). The k nearest neighbors are just the k users with the highest similarity scores to user `user_id`.
  - Each item in the returned list is a tuple consisting of two elements: the first element is a user\_id -- call it `uid_j` -- and the second element is the similarity score of user `uid_j` to user `user_id` -- call this `sim_j`

Sample output from `nearestNeighbors` would look like:

```
[("kluver", 0.91), ("lange", 0.78), ("tveite", 0.74), ("terveen", 0.37)]
```

To compute the k nearest neighbors, you must first compute the similarity of `user_id` to all the other users (be sure not to compare `user_id` to him/herself!). There are several ways to proceed next, but an effective way is simply to build a list of (`uid_j`, `sim_j`) pairs as you compute the similarity scores, then sort this list by the similarity scores, and finally return a list consisting of the first k items.

## A-Level

Your task at the A-level is to predict a rating for a user `U` and item `I`. We will assume you've already written and have access to the `nearestNeighbor` algorithm. Write the following function:

- `predict(item, k_nearest_neighbors, all_user_ratings)`
  - `item` is the item for which you want to make a prediction (for a given user).
  - `k_nearest_neighbors` is a list of tuples that is the output of `nearestNeighbors`.
  - `all_user_ratings` is the output of `readRatings`.
  - The output should be a float between 0.5 and 5.0.

**IMPORTANT CLARIFICATION:** The neighbors of some users for some  $k$  will not have ratings for every item. If the passed neighbors don't have ratings for the passed item return a prediction of 0.

Here is the idea for your solution:

- Remove any of the  $k$ -nearest-neighbors who don't have a rating for item.
- Compute a weight for each of the neighbors that determines how much each neighbor's opinion of item should influence the prediction for user  $U$ . Do this as follows:
  - Sum the similarities (to user  $U$ ) of each of the  $k$  neighbors. Call this  $S$ .  $S = \text{sum}(\text{sim}[i])$
- The prediction of how much user  $U$  will like item now is simply:
  - $\text{sum}(\text{sim}[i] * \text{rating}[i]) / S$ , where  $\text{rating}[i]$  is user  $i$ 's rating of item.

## Extra Credit

Your task is to build on what you've done at the previous level to write a function to recommend items to a target user  $U$ .

- `recommend(user_id, all_user_ratings)`
- Return a list of pairs (tuples) whose first elements are movies and whose second elements are the predicted ratings for the movie by user `user_id`, based on the data in `all_user_ratings` (which as usual is the output of `readRatings`)

Here's the idea:

- Find the  $k$ -nearest neighbors ( **$k=10$** )
- Find items that **any** of the neighbors have rated highly that user `user_id` has not rated. Define "rated highly" as ratings of 4.0, 4.5, and 5.0.
- This will give you a list of items (remember to remove duplicates).
- Now, compute a prediction for each of the items in the list (see A-level), and build the list of pairs: [(movie-1, prediction-1), ..., (movie-n, prediction-n)]
- Sort the list by prediction scores: you've got the recommendations!